

TAGA-LLM: Topology-Adaptive Graph Attention Informed Large Language Model for Multi-UAV Cooperative Search and Attack Mission Planning

Xinrui Guo¹, Jiang Zhao¹, Yingxun Wang¹, and Pei Chi¹

School of Automation Science and Electrical Engineering, Beihang University, Beijing
100191, China
guoxr@buaa.edu.cn, zhaojiang@buaa.edu.cn, yxwang@buaa.edu.cn,
peichi@buaa.edu.cn

Abstract. Cooperative search and attack missions for unmanned aerial vehicle (UAV) swarms require real-time coordination under dynamic communication topologies, yet existing approaches either rely on hand-crafted heuristics that lack adaptability or employ multi-agent reinforcement learning (MARL) policies that offer no interpretability. We propose **TAGA-LLM**, a two-stage framework that first trains a Topology-Adaptive Graph Attention enhanced Multi-Agent PPO (TAGA-MAPPO) expert policy, then distills its decision-making capability into a Qwen-4B large language model via LoRA fine-tuning with chain-of-thought supervision. The TAGA mechanism augments standard graph attention with three learnable, physics-informed factors—distance decay, task awareness, and information freshness—to model realistic communication constraints. Through expert trajectory collection, counterfactual data augmentation, and a topology-aware prompt encoding scheme, the distilled LLM achieves instruction-controllable and interpretable mission planning while preserving cooperative performance. We evaluate the framework on a 100×100 grid battlefield with 10 UAVs, 5 targets, and 7 threat zones, assessing both cooperative task performance and the quality of chain-of-thought reasoning across five operational directives.

Keywords: multi-UAV cooperation, graph attention network, large language model, multi-agent reinforcement learning, knowledge distillation

1 Introduction

Multi-UAV cooperative search and attack missions require a swarm of UAVs to autonomously explore unknown environments, discover hidden targets, and coordinate strikes under stringent communication constraints and threat zones [16]. The core challenge is achieving effective inter-agent coordination when each UAV possesses only a partial, local view and can communicate only within a limited radius.

Multi-Agent Reinforcement Learning (MARL) offers a data-driven approach to learning cooperative policies [11,22]. Graph-based communication mechanisms such as CommNet [18] and DGN [8] enable learned message passing, yet standard Graph Attention Networks (GATs) [19] treat all links equally, ignoring that wireless channel quality degrades with distance, co-modal agents share higher mutual information, and stale messages carry increasing uncertainty. Furthermore, MARL policies are inherently opaque. Large Language Models (LLMs) [2] have shown strong reasoning capabilities [21,17], but directly applying them to multi-UAV coordination is hindered by the lack of domain-specific cooperative strategies and the difficulty of encoding topology-dependent states into prompts.

We propose **TAGA-LLM**, a two-stage framework that combines MARL’s cooperative learning strength with LLM interpretability. The key insight is that MARL discovers effective cooperative strategies but locks them inside opaque neural networks, while LLMs excel at structured reasoning and natural language explanation but lack domain-specific cooperative knowledge; distillation bridges this complementary gap. In *Stage I*, we introduce the *Topology-Adaptive Graph Attention* (TAGA) mechanism with three learnable factors—distance decay, task awareness, and information freshness—integrated into a hierarchical MAPPO framework with curriculum learning. In *Stage II*, we distill the TAGA-MAPPO policy into a Qwen-4B LLM via LoRA fine-tuning with topology-aware prompt encoding, CoT supervision, and counterfactual data augmentation.

The primary contributions of this paper are summarized as follows: (a) a **TAGA mechanism** with three physics-informed adaptive factors for communication-aware multi-agent coordination; (b) a **two-stage MARL-to-LLM distillation pipeline** that yields interpretable and instruction-controllable mission planning; and (c) a **hierarchical policy architecture** trained with MAPPO and a three-stage curriculum.

The rest of this paper is organized as follows. In Sect. 2, we review related work. In Sect. 3, we state the Dec-POMDP formulation of the mission. In Sect. 4, we present the TAGA-LLM framework, including TAGA, MAPPO training, and distillation. In Sect. 5, we report simulation and evaluation plans. Finally, Sect. 6 concludes the paper.

2 Related Work

We briefly review multi-agent communication learning, LLM-based decision systems, policy distillation, and multi-UAV planning, and position TAGA-LLM relative to these lines of work.

Communication in MARL. CommNet [18] broadcasts mean-pooled messages; TarMAC [4] adds attention-based targeting; DGN [8] uses graph convolutions but assumes a static adjacency matrix. QMIX [13] and MAPPO [22] address credit assignment without modeling physical communication constraints. Our TAGA mechanism injects three physics-informed factors directly into the attention computation, treating the graph as a dynamic, quality-variable channel.

LLMs for decision-making. ReAct [21] interleaves reasoning with actions; Reflexion [17] adds self-reflection; Voyager [20] generates executable code; Say-Can [1] grounds LLM plans in affordances. These approaches either operate in single-agent settings or rely on pre-defined protocols, and none encodes topology-dependent multi-agent states into LLM prompts while preserving MARL-learned cooperative strategies.

Knowledge distillation for policy transfer. Policy distillation compresses a teacher policy into a smaller student [14]. Recent efforts explore distilling RL-learned behaviors into language models for interpretability [2, 10], yet existing pipelines do not account for graph-structured multi-agent communication or topology-dependent state encoding. Our Stage II distillation combines LoRA-based parameter-efficient fine-tuning [7] with topology-aware prompt encoding and chain-of-thought supervision to preserve cooperative strategies in the distilled LLM.

Multi-UAV planning. Classical methods include potential fields [6], auction-based allocation [3], and metaheuristics [5]. MARL-based approaches such as MADDPG [11] and HAPPO [9] train end-to-end policies, but for homogeneous swarms, parameter-sharing MAPPO is more efficient [22]. Our work builds on MAPPO with TAGA for communication-aware coordination and bridges the MARL-LLM gap through policy distillation.

3 Problem Statement

We model the mission as a Decentralized Partially Observable Markov Decision Process (Dec-POMDP) $\langle \mathcal{N}, \mathcal{S}, \{\mathcal{O}_i\}_{i=1}^N, \{\mathcal{A}_i\}_{i=1}^N, \mathcal{T}, \mathcal{R}, \gamma \rangle$ with $N = 10$ homogeneous UAVs on a 100×100 grid containing $M = 5$ hidden targets and $K = 7$ threat zones.

Observation. Each UAV i observes $o_i = (o_i^{\text{grid}}, o_i^{\text{self}})$: a five-channel 11×11 binary map (obstacles, explored cells, friendly UAVs, targets, threats) and a 6-dim state vector (position, velocity, energy, mode).

Action. UAV i selects mode $m_i \in \{\text{search}, \text{attack}\}$, then a mode-conditioned action: 9 moves (8 directions + stay) in search mode, plus *fire* in attack mode. Invalid actions are masked.

Communication. A dynamic graph $\mathcal{G}_t = (\mathcal{N}, \mathcal{E}_t)$ with edge $(i, j) \in \mathcal{E}_t$ iff $d_{ij} \leq R_{\text{comm}} = 20$. For each pair (i, j) , we track the *message staleness* $\Delta t_{ij} \geq 0$, defined as the number of time steps since UAV j last successfully transmitted to UAV i (reset to 0 upon reception).




Fig. 1. The TAGA-LLM framework. Stage I: TAGA-MAPPO with topology-adaptive graph attention and curriculum learning. Stage II: distillation to Qwen-4B via LoRA fine-tuning with chain-of-thought supervision.

Reward. The shared reward combines search coverage ($w_s = 0.1$), target discovery ($w_d = 1.0$), destruction ($w_a = 5.0$), cooperative bonus ($w_c = 3.0$), threat/loss penalty ($w_p = -2.0$), and time cost ($w_t = -0.01$):

$$r_t = w_s r_t^{\text{cov}} + w_d r_t^{\text{disc}} + w_a r_t^{\text{atk}} + w_c r_t^{\text{coop}} + w_p r_t^{\text{pen}} + w_t. \quad (1)$$

We evaluate via five metrics: search coverage, target destroy rate, completion time, UAV survival rate, and cooperative attack ratio.

4 Proposed Method

This section details the TAGA-LLM framework. As shown in Fig. 1, Stage I (Secs. 4.1–4.3) trains a TAGA-MAPPO expert policy, and Stage II (Sec. 4.4) distills it into an instruction-controllable LLM.

4.1 Topology-Adaptive Graph Attention

Given node features $\mathbf{h}_i \in \mathbb{R}^{128}$ from the observation encoder (Sec. 4.2), the attention logit between UAV i and neighbor $j \in \mathcal{N}_i$ at head k is $e_{ij}^{(k)} = (\mathbf{W}_Q^{(k)} \mathbf{h}_i)^\top (\mathbf{W}_K^{(k)} \mathbf{h}_j) / \sqrt{d_k}$, where $\mathbf{W}_Q^{(k)}, \mathbf{W}_K^{(k)} \in \mathbb{R}^{d_k \times 128}$ project into the k -th head space. TAGA modulates these logits with three multiplicative factors (dropping head superscripts for clarity):

$$\alpha_{ij}^{\text{TAGA}} = \frac{\exp(e_{ij} \cdot f_d(d_{ij}) \cdot f_m(m_i, m_j) \cdot f_t(\Delta t_{ij}))}{\sum_{j' \in \mathcal{N}_i} \exp(e_{ij'} \cdot f_d(d_{ij'}) \cdot f_m(m_i, m_{j'}) \cdot f_t(\Delta t_{ij'}))}, \quad (2)$$

where $\mathcal{N}_i = \{j : (i, j) \in \mathcal{E}_t\}$ denotes the communication neighbors of UAV i , and the three modulation factors are: (i) *distance decay* $f_d(d_{ij}) = \exp(-\lambda_d \cdot d_{ij}/R_{\text{comm}})$, modeling signal attenuation with distance (the indicator $\mathbb{I}(d_{ij} \leq$

R_{comm}) is implicit in the edge set $\mathcal{E}_t$); (ii) *task awareness* $f_m(m_i, m_j) = 1 + \lambda_m \cdot \mathbb{I}(m_i = m_j)$, amplifying attention between co-modal agents; (iii) *information freshness* $f_t(\Delta t_{ij}) = \exp(-\lambda_t \cdot \Delta t_{ij})$, discounting messages by their staleness Δt_{ij} (cf. Sec. 3). The coefficients $\lambda_d, \lambda_m, \lambda_t \in \mathbb{R}_{>0}$ are learnable scalars (initialized to 1.0, 1.0, 0.1, respectively). The multi-head aggregated output is:

$$\tilde{\mathbf{h}}_i = \text{LN} \left(\mathbf{W}_O \parallel_{k=1}^H \left(\sum_{j \in \mathcal{N}_i} \alpha_{ij}^{\text{TAGA},(k)} \mathbf{W}_V^{(k)} \mathbf{h}_j \right) + \mathbf{W}_R \mathbf{h}_i \right), \quad (3)$$

where $\parallel$ denotes concatenation across H heads and $\mathbf{W}_R$ is a residual projection handling dimension mismatch. We use 2 layers, $H = 4$ heads, $d_k = 16$, output dim 64.

4.2 Observation Encoder and Hierarchical Policy

The observation encoder has two branches: a CNN ($\text{Conv}_{5 \times 5}^{5 \times 5} \rightarrow \text{Conv}_{32 \times 32}^{5 \times 5} \rightarrow \text{AvgPool} \rightarrow \mathbb{R}^{128}$) for the $5 \times 11 \times 11$ grid, and an MLP ($6 \rightarrow 32 \rightarrow 64$) for the self-state. The concatenated features are projected to $\mathbf{h}_i \in \mathbb{R}^{128}$, which serves as input to the TAGA layers.

After TAGA, the enriched feature $\tilde{\mathbf{h}}_i \in \mathbb{R}^{64}$ is concatenated with $\mathbf{h}_i$ to form $\mathbf{f}_i = [\tilde{\mathbf{h}}_i; \mathbf{h}_i] \in \mathbb{R}^{192}$, which is fed into three policy heads: a *mode head* ($\mathbb{R}^{192} \rightarrow \mathbb{R}^2$: search/attack), a *search head* ($\mathbb{R}^{192} \rightarrow \mathbb{R}^9$: 8 directions + stay), and an *attack head* ($\mathbb{R}^{192} \rightarrow \mathbb{R}^{10}$: 9 moves + fire), all with hidden dimension 128. The joint action probability factorizes as $\pi(m_i, a_i | o_i) = \pi^m(m_i | \mathbf{f}_i) \cdot \pi^{m_i}(a_i | \mathbf{f}_i)$. Invalid actions are masked before the softmax normalization.

4.3 MAPPO Training with Curriculum Learning

We train with Centralized Training with Decentralized Execution (CTDE): a centralized critic receives mean-pooled TAGA features $\bar{\mathbf{h}} = \frac{1}{N} \sum_{i=1}^N \tilde{\mathbf{h}}_i$ and outputs a scalar state value $V_\phi(\bar{\mathbf{h}})$. Let $r_t(\theta) = \pi_\theta(a_t | o_t) / \pi_{\theta_{\text{old}}}(a_t | o_t)$ denote the importance sampling ratio. The clipped PPO objective is:

$$\mathcal{L}_{\text{PPO}} = -\mathbb{E}_t \left[\min(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1-\epsilon, 1+\epsilon) \hat{A}_t) \right] + c_v \mathcal{L}_{\text{val}} - c_e \mathcal{H}[\pi_\theta], \quad (4)$$

where $\hat{A}_t$ is the GAE advantage estimate [15] with $\gamma=0.99$ and $\lambda_{\text{GAE}}=0.95$, $\epsilon=0.2$ is the clipping threshold, $\mathcal{L}_{\text{val}} = (V_\phi(\mathbf{h}_t) - R_t)^2$ is the value loss, $\mathcal{H}[\pi_\theta]$ is the policy entropy bonus, and $c_v=0.5$, $c_e=0.01$ are balancing coefficients. Rewards are normalized with running mean and variance. Training uses a three-stage curriculum: easy (3 UAVs, 50^2 grid) $\rightarrow$ medium (6 UAVs, 75^2) $\rightarrow$ hard (10 UAVs, 100^2), advancing when the 20-episode moving average of search coverage exceeds a stage-specific threshold.

4.4 Stage II: TAGA-LLM Distillation

The trained TAGA-MAPPO policy achieves strong performance but remains a black box. Stage II distills it into an LLM for interpretability and instruction control.

Data collection. We deploy the converged TAGA-MAPPO policy for 3,000 episodes across three scenario configurations (varying target/threat layouts), converting each transition into a natural-language scene description paired with the expert action and a chain-of-thought rationale. To improve instruction controllability, we generate 15,000 counterfactual samples by re-labeling each scene under all five directives (BAL/SCH/ATK/EVD/COO). After filtering the top-75% by episode return, the final dataset contains $\sim 45\text{K}$ training samples.

Prompt encoding. Each sample is a structured prompt comprising four blocks: (a) self-state (position, velocity, energy, mode), (b) neighbor list sorted by ascending distance, where each entry is annotated with TAGA-derived signal quality (H/M/L, mapped from f_d) and freshness (f_t), (c) visible target list with coordinates and status, and (d) an operational instruction (BAL/SCH/ATK/EVD/COO). Priority-based truncation (max 256 tokens) removes the most distant neighbors first to stay within the context window.

Training. We fine-tune Qwen-4B [12] with LoRA [7] ($r = 16$, $\alpha = 32$, dropout 0.05) on all attention/FFN projections ($\sim 6.3\text{M}$ trainable params). Each sample includes CoT reasoning in `<think>` tags and a structured action in `<act>` tags. Let $\mathcal{M} \subset \{1, \dots, T\}$ denote the set of token positions within the `<think>` and `<act>` spans. The supervised fine-tuning loss restricts cross-entropy to these positions:

$$\mathcal{L}_{\text{SFT}} = -\frac{1}{|\mathcal{M}|} \sum_{t \in \mathcal{M}} \log p_{\theta}(x_t \mid x_{<t}). \quad (5)$$

Training completes in ~ 37 min on a single A100 with bf16 and gradient checkpointing. At inference, greedy decoding with 4-bit quantization achieves ~ 70 ms per step.

The complete training procedure is summarized in Algorithm 1.

5 Experiments and Simulation Results

5.1 Simulation Setup

To be completed: environment configuration, baseline descriptions (Random, Greedy, ACO, MAPPO, TAGA-MAPPO), evaluation protocol (30 independent runs per method, Welch’s t -test at $\alpha = 0.05$), and hardware specifications.

5.2 Simulation Results

To be completed: comparison of TAGA-LLM against baselines across five metrics (search coverage, target destroy rate, completion time, UAV survival rate, cooperative attack ratio). Results will be presented in a summary table with mean $\pm$ std and statistical significance markers.

Algorithm 1 Two-stage TAGA-LLM training and distillation procedure.

Require: Environment config $\mathcal{C}$, pre-trained Qwen-4B, curriculum thresholds $\{\tau_1, \tau_2\}$

Ensure: Distilled TAGA-LLM agent

- 1: **Stage I — TAGA-MAPPO Training**
 - 2: Initialize actor π_θ , critic V_ϕ , TAGA params $\lambda_d=1.0$, $\lambda_m=1.0$, $\lambda_t=0.1$
 - 3: Set curriculum stage $\leftarrow$ EASY (3 UAVs, 50^2 grid)
 - 4: **for** update = $1, \dots, T_{\max}$ **do**
 - 5: Collect rollout episodes with current π_θ
 - 6: Encode observations $\rightarrow$ TAGA features $\tilde{\mathbf{h}}_i$ (Eq. 2-3)
 - 7: Compute actions via hierarchical policy (Sec. 4.2)
 - 8: Update θ, ϕ via clipped PPO (Eq. 4)
 - 9: **if** 20-episode moving avg. coverage $> \tau_{\text{stage}}$ **then**
 - 10: Advance curriculum stage
 - 11: **end if**
 - 12: **end for**
 - 13: **Stage II — LLM Distillation**
 - 14: Deploy converged π_θ^* for 3K episodes $\rightarrow$ expert trajectories
 - 15: Convert to natural-language scenes; generate 15K counterfactual samples
 - 16: Filter top-75% by reward $\rightarrow$ dataset $\mathcal{D}$ ($\sim 45\text{K}$ samples)
 - 17: Attach LoRA ($r=16$, $\alpha=32$) to Qwen-4B
 - 18: Fine-tune on $\mathcal{D}$ with masked SFT loss (Eq. 5)
 - 19: **return** TAGA-LLM agent
-

5.3 Instruction Controllability

To be completed: performance under five operational directives (BAL, SCH, ATK, EVD, COO), demonstrating that TAGA-LLM responds meaningfully to different instructions.

5.4 Ablation Studies

To be completed: ablation on (i) TAGA-aware prompt encoding vs. flat prompt, (ii) CoT supervision vs. action-only supervision, (iii) counterfactual data augmentation, and (iv) individual TAGA factors (λ_d , λ_m , λ_t).

5.5 Scalability and Robustness

To be completed: scalability analysis with 5/10/15/20 UAVs, robustness under 10%/20%/30% communication dropout, and inference latency benchmarks.

6 Conclusion

This paper presented TAGA-LLM, a two-stage framework that bridges MARL cooperative performance with LLM interpretability for multi-UAV mission planning. The TAGA mechanism introduces three learnable, physics-informed attention factors—distance decay f_d , task awareness f_m , and information fresh-

ness f_t —that model realistic communication constraints beyond static-graph assumptions. The MARL-to-LLM distillation pipeline, combining topology-aware prompt encoding, chain-of-thought supervision, and counterfactual augmentation, transfers cooperative strategies into a Qwen-4B model that provides transparent reasoning, instruction controllability under the five mission directives (BAL-COO), and efficient edge deployment (~ 70 ms/step with 4-bit quantization).

Limitations. The current framework assumes homogeneous UAVs and a simplified grid-based environment. The distillation quality depends on the coverage of the expert trajectory dataset, and the 256-token prompt truncation may lose information in dense scenarios. Additionally, the chain-of-thought supervision relies on template-based rationale generation rather than human annotations.

Future work. Promising directions include extending to heterogeneous swarms with role-specific policies, incorporating online adaptation for non-stationary environments, and validating on real-world UAV testbeds with physical communication constraints.

References

1. Ahn, M., Brohan, A., Brown, N., Chebotar, Y., Cortes, O., David, B., Finn, C., Fu, C., Guber, K., Gopalakrishnan, K., et al.: Do as I can, not as I say: Grounding language in robotic affordances. In: Conference on Robot Learning. pp. 287–318 (2022)
2. Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J.D., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al.: Language models are few-shot learners. *Advances in Neural Information Processing Systems* 33, 1877–1901 (2020)
3. Choi, H.L., Brunet, L., How, J.P.: Consensus-based decentralized auctions for robust task allocation. *IEEE Transactions on Robotics* 25(4), 912–926 (2009)
4. Das, A., Gervet, T., Romoff, J., Batra, D., Parikh, D., Rabbat, M., Pineau, J.: TarMAC: Targeted multi-agent communication. In: International Conference on Machine Learning. pp. 1538–1546 (2019)
5. Dorigo, M., Birattari, M., Stutzle, T.: Ant colony optimization. *IEEE Computational Intelligence Magazine* 1(4), 28–39 (2006)
6. Ge, S.S., Cui, Y.J.: Dynamic motion planning for mobile robots using potential field method. *Autonomous Robots* 13(3), 207–222 (2002)
7. Hu, E.J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., Chen, W.: LoRA: Low-rank adaptation of large language models. In: International Conference on Learning Representations (2022)
8. Jiang, J., Dun, C., Huang, T., Lu, Z.: Graph convolutional reinforcement learning. In: International Conference on Learning Representations (2020)
9. Kuba, J.G., Chen, R., Wen, M., Wen, Y., Sun, F., Wang, J., Yang, Y.: Trust region policy optimisation in multi-agent reinforcement learning. In: International Conference on Learning Representations (2022)
10. Li, G., Hammoud, H.A.A.K., Itani, H., Khizbullin, D., Ghanem, B.: CAMEL: Communicative agents for “mind” exploration of large language model society. *Advances in Neural Information Processing Systems* 36 (2023)

11. Lowe, R., Wu, Y.I., Tamar, A., Harb, J., Abbeel, P., Mordatch, I.: Multi-agent actor-critic for mixed cooperative-competitive environments. In: *Advances in Neural Information Processing Systems*. vol. 30 (2017)
12. Qwen Team: Qwen2 technical report. arXiv preprint arXiv:2407.10671 (2024)
13. Rashid, T., Samvelyan, M., Schroeder de Witt, C., Farquhar, G., Foerster, J., Whiteson, S.: QMIX: Monotonic value function factorisation for deep multi-agent reinforcement learning. In: *International Conference on Machine Learning*. pp. 4295–4304 (2018)
14. Rusu, A.A., Colmenarejo, S.G., Gulcehre, C., Desjardins, G., Kirkpatrick, J., Pascanu, R., Mnih, V., Kavukcuoglu, K., Hadsell, R.: Policy distillation. *International Conference on Learning Representations* (2016)
15. Schulman, J., Moritz, P., Levine, S., Jordan, M., Abbeel, P.: High-dimensional continuous control using generalized advantage estimation. *International Conference on Learning Representations* (2016)
16. Shakhathreh, H., Sawalmeh, A.H., Al-Fuqaha, A., Dou, Z., Almaita, E., Khalil, I., Othman, N.S., Khreishah, A., Guizani, M.: Unmanned aerial vehicles (UAVs): A survey on civil applications and key research challenges. *IEEE Access* 7, 48572–48634 (2019)
17. Shinn, N., Cassano, F., Gopinath, A., Narasimhan, K., Yao, S.: Reflexion: Language agents with verbal reinforcement learning. In: *Advances in Neural Information Processing Systems*. vol. 36 (2023)
18. Sukhbaatar, S., Szlam, A., Fergus, R.: Learning multiagent communication with backpropagation. In: *Advances in Neural Information Processing Systems*. vol. 29 (2016)
19. Veličković, P., Cucurull, G., Casanova, A., Romero, A., Liò, P., Bengio, Y.: Graph attention networks. In: *International Conference on Learning Representations* (2018)
20. Wang, G., Xie, Y., Jiang, Y., Mandlekar, A., Xiao, C., Zhu, Y., Fan, L., Anandkumar, A.: Voyager: An open-ended embodied agent with large language models. arXiv preprint arXiv:2305.16291 (2023)
21. Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., Cao, Y.: ReAct: Synergizing reasoning and acting in language models. In: *International Conference on Learning Representations* (2023)
22. Yu, C., Velu, A., Vinitisky, E., Gao, J., Wang, Y., Baez, A., Bhatt, S., Finn, C., Wu, Y.: The surprising effectiveness of PPO in cooperative multi-agent games. In: *Advances in Neural Information Processing Systems*. vol. 35 (2022)